

Project - End-to-End Analysis Case Study

Data Cleaning

1. Load and Inspect the Dataset

```
In [39]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [40]: # Load the dataset  
df = pd.read_csv('Sample data (1).csv')  
df
```

Out[40]:

	Segment	Country	Product	Discount Band	Units Sold	Manufacturing Price	Sale Price	Gross Sales
0	Government	Germany	Carretera	NaN	1513.00	3	350	529550.00
1	Government	Germany	Paseo	NaN	1006.00	10	350	352100.00
2	Government	Canada	Paseo	NaN	1725.00	10	350	603750.00
3	Government	Germany	Paseo	NaN	1513.00	10	350	529550.00
4	Government	Germany	Velo	NaN	1006.00	120	350	352100.00
...	...	...	...	...	...	...	...	...
695	Midmarket	Canada	Paseo	High	1614.00	10	15	24210.00
696	Midmarket	Canada	Paseo	High	2559.00	10	15	38385.00
697	Enterprise	Germany	Paseo	High	1085.00	10	125	135625.00
698	Midmarket	Germany	Paseo	High	1175.00	10	15	17625.00
699	Channel Partners	United States of America	Paseo	High	914.00	10	12	10968.00

700 rows × 16 columns



In [41]: df.info()

```
<class 'pandas.DataFrame'>
RangeIndex: 700 entries, 0 to 699
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Segment                700 non-null   str
1   Country                700 non-null   str
2   Product                700 non-null   str
3   Discount Band          647 non-null   str
4   Units Sold             700 non-null   float64
5   Manufacturing Price     700 non-null   int64
6   Sale Price             700 non-null   int64
7   Gross Sales            700 non-null   float64
8   Discounts              700 non-null   float64
9   Sales                  700 non-null   float64
10  COGS                   700 non-null   float64
11  Profit                 700 non-null   float64
12  Date                   700 non-null   str
13  Month Number           700 non-null   int64
14  Month Name             700 non-null   str
15  Year                   700 non-null   int64
dtypes: float64(6), int64(4), str(6)
memory usage: 118.9 KB
```

2. View the Actual Duplicate Rows

To see the rows that are duplicated so you can inspect them:

```
In [42]: # Keep='false' shows all instances of the duplicated rows
duplicate_rows = df[df.duplicated(keep=False)]

display(duplicate_rows)
```

Segment	Country	Product	Discount Band	Units Sold	Manufacturing Price	Sale Price	Gross Sales	Discounts	Sale

a) Is null ?

```
In [43]: df.isnull().sum()
```

```
Out[43]: Segment      0
Country      0
Product      0
Discount Band 53
Units Sold   0
Manufacturing Price 0
Sale Price   0
Gross Sales  0
Discounts    0
Sales        0
COGS         0
Profit       0
Date         0
Month Number 0
Month Name   0
Year         0
dtype: int64
```

```
In [44]: df.head()
```

```
Out[44]:
```

	Segment	Country	Product	Discount Band	Units Sold	Manufacturing Price	Sale Price	Gross Sales	Di
0	Government	Germany	Carretera	NaN	1513.00	3	350	529550.00	
1	Government	Germany	Paseo	NaN	1006.00	10	350	352100.00	
2	Government	Canada	Paseo	NaN	1725.00	10	350	603750.00	
3	Government	Germany	Paseo	NaN	1513.00	10	350	529550.00	
4	Government	Germany	Velo	NaN	1006.00	120	350	352100.00	

b) Cecking each column for mixed value types

```
In [45]: # 1. Check the general data type assigned to each column
print(df.dtypes)

print("\n-----\n")

# 2. Check for mixed Python types within the same column
for col in df.columns:
    unique_types = df[col].apply(type).unique()
    print(f"Column '{col}': {[t.__name__ for t in unique_types]}")
```

```

Segment          str
Country          str
Product          str
Discount Band    str
Units Sold       float64
Manufacturing Price int64
Sale Price       int64
Gross Sales      float64
Discounts        float64
Sales            float64
COGS             float64
Profit           float64
Date             str
Month Number     int64
Month Name       str
Year            int64
dtype: object

```

```

-----

Column 'Segment': ['str']
Column 'Country': ['str']
Column 'Product': ['str']
Column 'Discount Band': ['float', 'str']
Column 'Units Sold': ['float']
Column 'Manufacturing Price': ['int']
Column 'Sale Price': ['int']
Column 'Gross Sales': ['float']
Column 'Discounts': ['float']
Column 'Sales': ['float']
Column 'COGS': ['float']
Column 'Profit': ['float']
Column 'Date': ['str']
Column 'Month Number': ['int']
Column 'Month Name': ['str']
Column 'Year': ['int']

```

Only Column 'Discount Band' has mixed type ['float', 'str']

c) Whats is in 'Discount Band' column?

```
In [46]: df['Discount Band'].unique()
```

```
Out[46]: <ArrowStringArray>
[nan, 'Low', 'Medium', 'High']
Length: 4, dtype: str
```

d) Get all null (NaN) rows

```
In [47]: # Filter rows where 'Discount Band' is NaN (null)
nan_discount_rows = df[df['Discount Band'].isna()] # isna() means NaN or null

# View the count and first few rows
```

```
print(f"Total rows with NaN Discount Band: {len(nan_discount_rows)}")
nan_discount_rows.head()
```

Total rows with NaN Discount Band: 53

Out[47]:

	Segment	Country	Product	Discount Band	Units Sold	Manufacturing Price	Sale Price	Gross Sales	Di
0	Government	Germany	Carretera	NaN	1513.00	3	350	529550.00	
1	Government	Germany	Paseo	NaN	1006.00	10	350	352100.00	
2	Government	Canada	Paseo	NaN	1725.00	10	350	603750.00	
3	Government	Germany	Paseo	NaN	1513.00	10	350	529550.00	
4	Government	Germany	Velo	NaN	1006.00	120	350	352100.00	

2. Clean the Dataset

```
In [48]: # Clean column names (strip trailing/leading spaces)
df.columns = df.columns.str.strip()

# Handle missing values in 'Discount Band' (replace NaN with 'None')
df['Discount Band'] = df['Discount Band'].fillna('None')

# Convert 'Date' column to datetime format
# Explicitly specify the date format to avoid the UserWarning
df['Date'] = pd.to_datetime(df['Date'], format='%d-%b-%y')

# Verify cleaning results
print(df.info())
print("\nUnique values in Discount Band:", df['Discount Band'].unique())
```

```

<class 'pandas.DataFrame'>
RangeIndex: 700 entries, 0 to 699
Data columns (total 16 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   Segment               700 non-null   str    
 1   Country               700 non-null   str    
 2   Product               700 non-null   str    
 3   Discount Band         700 non-null   str    
 4   Units Sold            700 non-null   float64
 5   Manufacturing Price    700 non-null   int64  
 6   Sale Price            700 non-null   int64  
 7   Gross Sales           700 non-null   float64
 8   Discounts              700 non-null   float64
 9   Sales                 700 non-null   float64
10  COGS                  700 non-null   float64
11  Profit                700 non-null   float64
12  Date                  700 non-null   datetime64[us]
13  Month Number          700 non-null   int64  
14  Month Name            700 non-null   str    
15  Year                  700 non-null   int64  
dtypes: datetime64[us](1), float64(6), int64(4), str(5)
memory usage: 113.5 KB
None

Unique values in Discount Band: <ArrowStringArray>
['None', 'Low', 'Medium', 'High']
Length: 4, dtype: str

```

a) Checked again for null after cleaning done

```
In [49]: df.isnull().sum()
```

```

Out[49]: Segment                0
Country                0
Product                0
Discount Band          0
Units Sold             0
Manufacturing Price    0
Sale Price             0
Gross Sales            0
Discounts              0
Sales                  0
COGS                   0
Profit                 0
Date                   0
Month Number           0
Month Name             0
Year                   0
dtype: int64

```

3. Save the cleaned dataset to a new CSV file

```
In [38]: df.to_csv('cleaned_retail_sales.csv', index=False)
```

```
print("\nCleaned dataset successfully saved as 'cleaned_retail_sales.csv'")
```

Cleaned dataset successfully saved as 'cleaned_retail_sales.csv'

Exploratory Data Analysis (EDA)

```
In [18]: # Set pandas to show numbers normally (up to 2 decimal places) & not use e scientific
pd.set_option('display.float_format', lambda x: '%.2f' % x)
```

```
In [19]: # Summary statistics for numerical columns
print(df.describe())
```

	Units Sold	Manufacturing Price	Sale Price	Gross Sales	Discounts	\
count	700.00	700.00	700.00	700.00	700.00	
mean	1608.29	96.48	118.43	182759.43	13150.35	
min	200.00	3.00	7.00	1799.00	0.00	
25%	905.00	5.00	12.00	17391.75	800.32	
50%	1542.50	10.00	20.00	37980.00	2585.25	
75%	2229.12	250.00	300.00	279025.00	15956.34	
max	4492.50	260.00	350.00	1207500.00	149677.50	
std	867.43	108.60	136.78	254262.28	22962.93	

	Sales	COGS	Profit	Date	Month	Number	\
count	700.00	700.00	700.00	700		700.00	
mean	169609.07	145475.21	24133.86	2014-04-28 21:36:00		7.90	
min	1655.08	918.00	-40617.50	2013-09-01 00:00:00		1.00	
25%	15928.00	7490.00	2805.96	2013-12-24 06:00:00		5.75	
50%	35540.20	22506.25	9242.20	2014-05-16 12:00:00		9.00	
75%	261077.50	245607.50	22662.00	2014-09-08 12:00:00		10.25	
max	1159200.00	950625.00	262200.00	2014-12-01 00:00:00		12.00	
std	236726.35	203865.51	42760.63	NaN		3.38	

	Year
count	700.00
mean	2013.75
min	2013.00
25%	2013.75
50%	2014.00
75%	2014.00
max	2014.00
std	0.43

Findings : Core Statistical Findings

Sales & Profit Distribution: *The dataset contains 700 transactions spanning across 2013 and 2014. Mean sales per transaction are roughly 169,609 with an average profit of 24,133, though profitability features a wide range (min profit is -40,617 representing financial losses on specific transactions, up to a max profit of 262,200).*

Units Sold: *The average units sold per transaction is 1,608 items, spanning from a minimum of 200 up to 4,492 units.*

b) Group By Country with most sales on top

```
In [20]: # Total Profit and Sales by Country
country_performance = df.groupby('Country')[['Sales', 'Profit', 'Gross Sales', 'Units Sold']]
country_performance
```

```
Out[20]:
```

	Country	Sales	Profit	Gross Sales	Units Sold
4	United States of America	25029830.16	2995540.67	27269358.00	232627.50
0	Canada	24887654.89	3529228.89	26932163.50	247428.50
1	France	24354172.28	3781020.78	26081674.50	240931.00
2	Germany	23505340.82	3680388.82	24921467.50	201494.00
3	Mexico	20949352.11	2907523.11	22726935.00	203325.00

c) Group by Year with most sales on top

```
In [21]: # by Year
year_performance = df.groupby('Year')[['Sales', 'Profit', 'Gross Sales', 'Units Sold']]
year_performance
```

```
Out[21]:
```

	Year	Sales	Profit	Gross Sales	Units Sold
1	2014	92311094.75	13015237.75	99370811.50	861132.00
0	2013	26415255.51	3878464.51	28560787.00	264674.00

d) Group by Month with most sales on top

```
In [22]: # by Month
month_performance = df.groupby('Month Name')[['Sales', 'Profit', 'Gross Sales', 'Units Sold']]
month_performance
```

Out[22]:

	Month Name	Sales	Profit	Gross Sales	Units Sold
10	October	21671431.02	3439781.02	23142112.00	201104.00
2	December	17367228.98	2717329.98	18343293.00	155306.00
9	November	12651417.50	1370102.50	14115248.00	121131.00
11	September	10882697.27	1786735.27	11575053.00	107881.00
6	June	9518893.82	1473753.82	10268972.00	103302.00
5	July	8102920.18	923865.68	8833027.50	69349.00
3	February	7297531.39	1148547.39	7699201.00	55115.00
0	April	6964775.07	929984.57	7429392.50	78886.50
4	January	6607761.68	814028.68	7307403.50	67835.50
8	May	6210211.06	828640.06	6767911.00	51771.00
1	August	5864622.42	791066.42	6325959.00	60705.00
7	March	5586859.87	669866.87	6124026.00	53420.00

e) Group by Segment with most sales on top

In [23]:

```
# by Segment
segment_performance = df.groupby('Segment')[['Sales', 'Profit', 'Gross Sales', 'Units Sold']]
segment_performance
```

Out[23]:

	Segment	Sales	Profit	Gross Sales	Units Sold
2	Government	52504260.67	11388173.17	56403066.50	470673.50
4	Small Business	42427918.50	4143168.50	45941700.00	153139.00
1	Enterprise	19611694.38	-614545.62	21069000.00	168552.00
3	Midmarket	2381883.08	660103.07	2582670.00	172178.00
0	Channel Partners	1800593.64	1316803.14	1935162.00	161263.50

Segment Performance Breakdown

Government is the largest customer segment, driving over 52.5 million in total sales and generating more than \$11.3 million in net profit.

Small Business represents the second highest sales driver (42.4M), while Enterprise registers a net negative total profit (-\$614,545) across its transaction lines, highlighting an area of structural margin loss.

f) Group by Product with most sales on top

```
In [24]: # by Product
product_performance = df.groupby('Product')[['Sales', 'Profit', 'Gross Sales', 'Units Sold']]
product_performance
```

```
Out[24]:
```

	Product	Sales	Profit	Gross Sales	Units Sold
3	Paseo	33011143.95	4797437.95	35611662.00	338239.50
4	VTT	20511921.02	3034608.02	21968533.50	168783.00
5	Velo	18250059.46	2305992.46	19826768.50	162424.50
0	Amarilla	17747116.06	2814104.06	19037279.50	155315.00
2	Montana	15390801.88	2114754.88	16549834.50	154198.00
1	Carretera	13815307.88	1826804.89	14937520.50	146846.00

Product Performance Breakdown

Paseo is the dominant product line by a significant margin, pulling in over 33 million in sales and ~338,239 units sold.

VTT and Velo follow as second and third highest revenue generators, whereas Carretera registers the lowest aggregate sales among the 6 products.

Data Visualization

1. Total Sales by Country Bar Chart


```
In [25]: import matplotlib.pyplot as plt
import seaborn as sns

sns.set_theme(style="whitegrid")
plt.figure(figsize=(12, 5))

country_sales = df.groupby('Country')['Sales'].sum().reset_index()

# Fix applied: hue='Country' and Legend=False
sns.barplot(data=country_sales, x='Country', y='Sales', hue='Country', palette='magma')

plt.title('Total Sales by Country')
plt.xlabel('Country')
plt.ylabel('Total Sales')
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()
```



The United States of America leads total sales closely followed by Canada and France, while Mexico records the lowest overall performance.

Sales distribution across the top four countries remains remarkably balanced, with each generating between USD 23.5 million and USD 25.0 million.


2. Total Profit by Product Bar Chart

```
In [26]: plt.figure(figsize=(10, 5))

product_profit = df.groupby('Product')['Profit'].sum().reset_index()

# Fix applied: hue='Product' and Legend=False
sns.barplot(data=product_profit, x='Product', y='Profit', hue='Product', palette='v


plt.title('Total Profit by Product')
plt.xlabel('Product')
plt.ylabel('Total Profit')
plt.tight_layout()
plt.show()
```



Paseo generated the highest total profit by a wide margin, whereas Carretera recorded the lowest profitability among all products

3. Segment Sales Bar Chart


```
In [27]: plt.figure(figsize=(10, 5))
segment_sales = df.groupby('Segment')['Sales'].sum().reset_index()
sns.barplot(data=segment_sales, x='Segment', y='Sales', hue='Segment', palette='crest')
plt.title('Total Sales by Segment')
plt.xlabel('Segment')
plt.ylabel('Total Sales')
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()
```



Government and Small Business segments drive the vast majority of total sales, significantly outperforming Enterprise, Midmarket, and Channel Partners.

4. Month Sales Bar Chart (Chronological Order)


```
In [28]: plt.figure(figsize=(12, 5))
month_order = ['January', 'February', 'March', 'April', 'May', 'June', 'July', 'Aug
month_sales = df.groupby('Month Name')['Sales'].sum().reindex(month_order).reset_in
sns.barplot(data=month_sales, x='Month Name', y='Sales', hue='Month Name', palette=
plt.title('Total Sales by Month')
plt.xlabel('Month')
plt.ylabel('Total Sales')
plt.xticks(rotation=30)
plt.tight_layout()
plt.show()
```



Sales peak dramatically during October and December, while earlier months like March, May, and August experience lower overall sales volume.

5. Year Sales Bar Chart

```
In [29]: plt.figure(figsize=(8, 3))
year_sales = df.groupby('Year')['Sales'].sum().reset_index()
year_sales['Year'] = year_sales['Year'].astype(str)
sns.barplot(data=year_sales, x= 'Sales', y= 'Year', hue='Year', palette='rocket')
plt.title('Total Sales by Year')
plt.xlabel('Year')
plt.ylabel('Total Sales')
plt.tight_layout()
plt.show()
```



Total sales in 2014 expanded significantly compared to 2013, more than tripling the previous year's annual revenue.


6. Correlation Matrix Heatmap

```
In [30]: import numpy as np

plt.figure(figsize=(9, 7))
numeric_df = df.select_dtypes(include=[np.number])
```

```
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm', fmt='.2f', linewidths=0)

plt.title('Correlation Matrix of Dataset Numerical Features')
plt.tight_layout()
plt.show()
```



Key Findings from the Correlation Matrix:

Gross Sales & Sales (r = 0.998): There is a near-perfect positive correlation between Gross Sales and Sales.


Cost of Goods Sold (COGS) & Sales (r = 0.992): COGS scales tightly with Sales and Gross Sales, pointing to a consistent proportional cost baseline for products sold regardless of volume.

Sale Price & Gross Sales (r = 0.808): Higher base sale prices strongly drive up gross sales totals, showing that price point strategy heavily dictates aggregate revenue capacity compared to unit volume alone.

Profit & Sales (r = 0.805): A robust correlation confirms that higher sales volumes successfully translate into higher net profitability across the dataset's transactions.

7. Scatter Plot Code (Units Sold vs Sales)


```
In [33]: plt.figure(figsize=(10, 6))
sns.scatterplot(data=df, x='Units Sold', y='Sales', hue='Segment', palette='deep',
plt.title('Scatter Plot: Units Sold vs Sales by Segment')
plt.xlabel('Units Sold')
plt.ylabel('Sales')
plt.tight_layout()
plt.show()
```



Sales scale positively with units sold across all segments, featuring distinct pricing tiers that separate **high-volume segments like Government and Small Business from lower-tier ones**.

8. Box Plot Code (Profit Distribution by Product)

```
In [32]: plt.figure(figsize=(10, 6))
sns.boxplot(data=df, x='Product', y='Profit', hue='Product', palette='Set2')
plt.title('Box Plot: Profit Distribution by Product')
plt.xlabel('Product')
plt.ylabel('Profit')
plt.tight_layout()
plt.show()
```



Paseo exhibits the highest profit variability and peak outlier returns, whereas Carretera and Montana show more conservative and tightly clustered profit distributions.

```
In [ ]:
```